




GPT-2 (124M)



Llama-3.2-1B



Llama-3.2-3B



inference compute (relative to greedy, log scale)

○ likelihood: 1 greedy, 2 beam-4, 3 greedy+rp △ sampling: 4 T=0.7, 5 T=1.3, 6 nucleus
 □ contrastive: 7 DoLa, 8 DoLa+nuc ◇ aggregation: 9 SC k=5 dashed line: Pareto frontier